

# Operating System Concepts

---

## Agenda

- syscall internals
- process hierarchy
- orphan & zombie state
- thread vs process
- thread creation
- synchronization
- semaphore vs mutex
- deadlock

## syscall internals

- Refer yesterday's slides

## Process Creation

- System Calls
  - Windows: CreateProcess()
  - UNIX: fork()
  - BSD UNIX: fork(), vfork()
  - Linux: clone(), fork(), vfork()

## fork() syscall

- To execute certain task concurrently we can create a new process (using fork() on UNIX).
- fork() creates a new process by duplicating calling process.
- The new process is called as "child process", while calling process is called as "parent process".
- "child" process is exact duplicate of the "parent" process except few points pid, parent pid, etc.

- `pid = fork();`
  - On success, `fork()` returns `pid` of the child to the parent process and 0 to the child process.
  - On failure, `fork()` returns -1 to the parent.
- Even if child is copy of the parent process, after its creation it is independent of parent and both these processes will be scheduled separately by the scheduler.
- Based on CPU time given for each process, both processes will execute concurrently.

### `getpid()` vs `getppid()`

- `pid1 = getpid();` // returns `pid` of the current process
- `pid2 = getppid();` // returns `pid` of the parent of the current process

### When `fork()` will fail?

- When no new PCB can be allocated, then `fork()` will fail.
- Linux has max process limit for the system and the user. When try to create more processes, `fork()` fails.
- terminal> `cat /proc/sys/kernel/pid_max`

### Orphan process

- If parent of any process is terminated, that child process is known as orphan process.
- The ownership of such orphan process will be taken by "init" process.

### Zombie process

- If process is terminated before its parent process and parent process is not reading its exit status, then even if process's memory/resources is released, its PCB will be maintained. This state is known as "zombie state".
- To avoid zombie state parent process should read exit status of the child process. It can be done using `wait()` syscall.

### `wait()` syscall

- `ret = wait(&s);`
  - `arg1`: out param to get exit code of the process.
  - returns: `pid` of the child process whose exit code is collected.

- wait() performs 3 steps:
  - Pause execution parent until child process is terminated.
  - Read exit code from PCB of child process & return to parent process (via out param).
  - Release PCB of the child process.
- The exit status returned by the wait() contains exit status, reason of termination and other details.
- Few macros are provided to access details from the exit code.
  - WEXITSTATUS()

## waitpid() syscall

- This extended version of wait() in Linux.
- `ret = waitpid(child_pid, &s, flags);`
  - arg1: pid of the child for which parent should wait.
    - -1 means any child.
  - arg2: out param to get exit code of the process.
  - arg3: extra flags to define behaviour of waitpid().
- returns: pid of the child process whose exit code is collected.
  - -1: if error occurred.

## exec() syscall

- exec() syscall "loads a new program" in the calling process's memory (address space) and replaces the older (calling) one.
- If exec() succeed, it does not return (rather new program is executed).
- There are multiple functions in the family of exec():
  - `execl()`, `execvp()`, `execle()`,
  - `execv()`, `execvp()`, `execve()`, `execvpe()`
- exec() family multiple functions have different syntaxes but same functionality.

## Thread concept

- Threads are used to execute multiple tasks concurrently in the same program/process.
- Thread is a light-weight process.
  - For each thread new control block and stack is created. Other sections (text, data, heap, ...) are shared with the parent process.

- Inter-thread communication is much faster than inter-process communication.
- Context switch between two threads in the same process is faster.
- Thread stack is used to create function activation records of the functions called/executed by the thread.

## Process vs Thread

- In modern OS, process is a container holding resources required for execution, while thread is unit of execution/scheduling.
- Process holds resources like memory, open files, IPC (e.g. signal table, shared memory, pipe, etc.).
- PCB contains resources information like pid, exit status, open files, signals/ipc, memory info, etc.
- CPU time is allocated to the threads. Thread is unit of execution.
- TCB contains execution information like tid, scheduling info (priority, sched algo, time left, ...), Execution context, Kernel stack, etc.
- terminal> ps -e -o pid,nlwp,cmd
- terminal> ps -e -m -o pid,tid,nlwp

## main thread

- For each process one thread is created by default called as main thread.
- The main thread executes entry-point function of the process.
- The main thread use the process stack.
- When main thread is terminated, the process is terminated.
- When a process is terminated, all threads in the process are terminated.

## thread functions

- pthread\_create()
  - Create a new thread.
    - arg1: posix thread id (out param)
    - arg2: thread attributes -- NULL means default attributes
      - stack size
      - scheduling policy
      - priority
    - arg3: address of thread function

- `void* thread_function(void*);`
  - `arg1: void*` -- param to the thread function (can be of any type, array or struct).
  - `returns: void*` -- result of thread (can be of any type)
- `arg4: param to thread function`
- `returns: 0` on success.
- Join thread
  - The current thread wait for completion of given thread and will collect return value of that thread.
  - `pthread_join(th_id, &res);`
    - `arg1: given thread` (for which current thread is to blocked).
    - `arg2: address of result variable` (out param to collect result of the given thread)
- Thread termination
  - When thread function is completed, the thread is terminated.
  - To terminate current thread early use `pthread_exit()` function.
  - `pthread_exit(result);`
    - `arg1: result (void*)` of the current thread
- Thread cancellation
  - A thread in a process can request to cancel execution of another thread.
  - `pthread_cancel(tid)`
    - `arg1: id of the thread to be cancelled.`

## Threading models

- Threads created by thread libraries are used to execute functions in user program. They are called as "user threads".
- Threads created by the syscalls (or internally into the kernel) are scheduled by kernel scheduler. They are called as "kernel threads".
- User threads are dependent on the kernel threads. Their dependency/relation (managed by thread library) is called as "threading model".
- There are four threading models:
  - Many to One
  - Many to Many

- One To One
- Two Level Model
- Many to One
  - Many user threads depends on single kernel thread.
  - If one of the user thread is blocked, remaining user threads cannot function.
- Many to Many
  - Many user threads depend on equal or less number of kernel threads.
  - If one of the user thread is blocked, other user thread keep executing (based on remaining kernel threads).
- One To One
  - One user thread depends on one kernel thread.
- Two Level Model
  - OS/Thread library supports both one to one and many to many model

## Synchronization

- Multiple processes accessing same resource at the same time, is known as "race condition".
- When race condition occurs, resource may get corrupted (unexpected results).
- Peterson's problem, if two processes are trying to modify same variable at the same time, it can produce unexpected results.
- Code block to be executed by only one process at a time is referred as Critical section. If multiple processes execute the same code concurrently it may produce undesired results.
- To resolve race condition problem, one process can access resource at a time. This can be done using sync objects/primitives given by OS.
- OS Synchronization objects are:
  - Semaphore, Mutex

## Semaphore

- Semaphore is a sync primitive given by OS.
- Internally semaphore is a counter. On semaphore two operations are supported:
  - wait operation: dec op: P operation:
    - semaphore count is decremented by 1.
    - if  $cnt < 0$ , then calling process is blocked.
    - typically wait operation is performed before accessing the resource.

- signal operation: inc op: V operation:
  - semaphore count is incremented by 1.
  - if one or more processes are blocked on the semaphore, then one of the process will be resumed.
  - typically signal operation is performed after releasing the resource.
- Semaphore types
  - Counting Semaphore
    - Allow "n" number of processes to access resource at a time.
    - Or allow "n" resources to be allocated to the process.
  - Binary Semaphore
    - Allows only 1 process to access resource at a time or used as a flag/condition.

## Mutex

- Mutex is used to ensure that only one process can access the resource at a time.
- Functionally it is same as "binary semaphore".
- Mutex can be unlocked by the same process/thread, which had locked it.

## Semaphore vs Mutex

- S: Semaphore can be decremented by one process and incremented by same or another process.
- M: The process locking the mutex is owner of it. Only owner can unlock that mutex.
- S: Semaphore can be counting or binary.
- M: Mutex is like binary semaphore. Only two states: locked and unlocked.
- S: Semaphore can be used for counting, mutual exclusion or as a flag.
- M: Mutex can be used only for mutual exclusion.

## Deadlock

- Deadlock occurs when four conditions/characteristics hold true at the same time.
  - No preemption: A resource should not be released until task is completed.
  - Mutual exclusion: Resources is not sharable.
  - Hold & Wait: Process holds a resource and wait for another resource.
  - Circular wait: Process P1 holds a resource needed for P2, P2 holds a resource needed for P3 and P3 holds a resource needed for P1.

## Deadlock Prevention

- OS syscalls are designed so that at least one deadlock condition does not hold true.
- In UNIX multiple semaphore operations can be done at the same time.

## Deadlock Avoidance

- Processes declare the required resources in advanced, based on which OS decides whether resource should be given to the process or not.
- Algorithms used for this are:
  - Resource allocation graph: OS maintains graph of resources and processes. A cycle in graph indicate circular wait will occur. In this case OS can deny a resource to a process.
  - Banker's algorithm: A bank always manage its cash so that they can satisfy all customers.
  - Safe state algorithm: OS maintains statistics of number of resources and number processes. Based on stats it decides whether giving resource to a process is safe or not (using a formula):
    - $\text{Max num of resources required} < \text{Num of resources} + \text{Num of processes}$ 
      - If condition is true, deadlock will never occur.
      - If condition is false, deadlock may occur

## IPC overview

- A process cannot access of memory of another process directly. OS provides IPC mechanisms so that processes can communicate with each other.
- IPC models
  - Shared memory model
    - Processes write/read from the memory region accessible to both the processes.
    - OS only provides access to the shared memory region.
  - Message passing model
    - Process send message to the OS and the other process receives message from the OS.
    - This is slower compared to shared memory model.
- Unix/Linux IPC mechanisms
  - Signals
  - Shared memory
  - Message queue



- Pipe
- Socket

## VI Editor

- Refer slides
- `sudo apt-get install vim`
- VI editor works in two modes
  - command mode
  - insert mode
- press `i` - to go into insert mode
- press `Esc` - to go into command mode
- VI editor commands:
  - `:w` - write/save into file
  - `:q` - quit vi editor
  - `:wq` - save and quit
  - `:y` - to copy current line
  - `yy` - to copy current line
  - `nyy` - copy `n` lines from current line
  - `:m,ny` - copy from `m`th line to `n`th line
  - `:d` - to cut current line
  - `dd` - to cut current line
  - `ndd` - cut `n` lines from current line
  - `:m,nd` - cut from `m`th line to `n`th line

- press p - to paste copied line on next line of current line
- To do setting in vi editor
  - create file into home directory as below:

```
vim ~/.vimrc
```

- add below content into it

```
set number  
set autoindent  
set tabstop=4  
set nowrap
```